

Provably Safe Controller Synthesis Using Safety Proofs as Building Blocks

Yanni Kouskoulas¹ Aurora Schmidt¹
Jean-Baptiste Jeannin² Daniel Genin¹ Jessica Lopez¹

International Conference in Software Engineering
Research and Innovation 2019

This work was supported by the Federal Aviation Administration (FAA)
Traffic-Alert & Collision Avoidance System (TCAS) Program Office (PO)
AJM-233: Volpe National Transportation Systems Center

Contract Nos. DTFAWA11C00074 and DTRT5715D30011

¹Johns Hopkins University Applied Physics Laboratory

²University of Michigan

OUTLINE

- ▶ Introduction
- ▶ Approach
- ▶ Controller Synthesis
- ▶ Implementation and Results
- ▶ Conclusions

OUTLINE

- ▶ **Introduction**
- ▶ Approach
- ▶ Controller Synthesis
- ▶ Implementation and Results
- ▶ Conclusions

PROBLEM

- ▶ Control software for complex, safety-critical robotic machinery needs safety guarantees
 - ▶ Primary controller based on machine learning approaches or neural networks may not have desired safety guarantees



Image from airlive.net



Image from Intuitive Surgical

MOTIVATING QUESTION

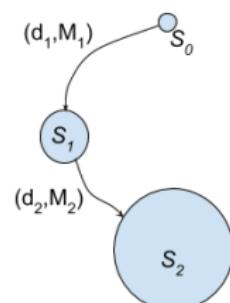
- ▶ How can we impose safety guarantees on cyber-physical systems based on machine learning techniques, or whose algorithms learn and evolve over time?

OUTLINE

- ▶ Introduction
- ▶ **Approach**
- ▶ Controller Synthesis
- ▶ Implementation and Results
- ▶ Conclusions

MODELING SYSTEM DYNAMICS

- ▶ Develop hybrid system models that represent sequences of discrete actions and continuous physics of the system
 - ▶ Differential equations with uncertain parameters
- ▶ Predictive, non-deterministic, continuous model represents ranges of future possibilities that encompass realistic but bounded behavior
- ▶ Schematic illustration of uncertainty growth in future projection of action plan.
 - ▶ Set S_0 corresponds to the initial known state of the actor. Sets S_1 and S_2 correspond to reachable sets at the end of actions (d_1, M_1) and (d_2, M_2) , respectively.



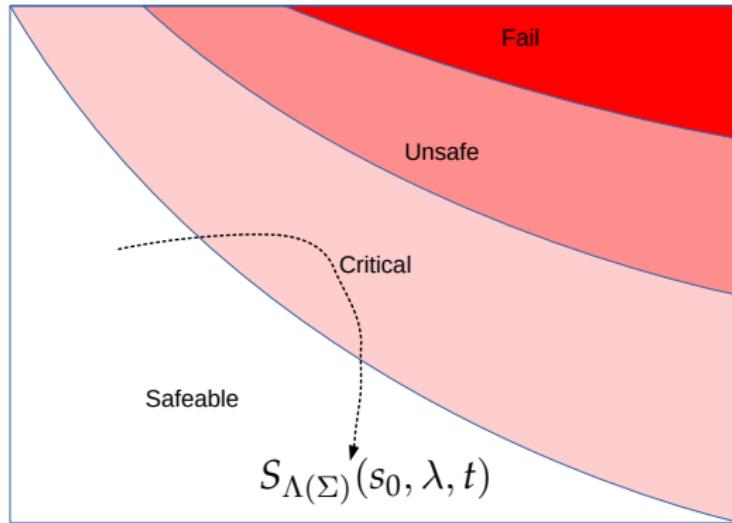
MAIN IDEA

- ▶ Improve the safety of computer-controlled robotic machinery by using safety proofs as building material for runtime safety controller
 - ▶ Construct system component from safety proofs
 - ▶ Invoke formally verified safety calculations at runtime
 - ▶ Confer safety guarantees provided by proofs on primary controller
 - ▶ What are minimal requirements for easy translation?

APPROACH

- ▶ Work with a dictionary of action sequences
- ▶ Identify where we are in the state space according to the availability of safety-preserving action sequences (SPAS)
 - ▶ A (SPAS) is one that is *guaranteed* to preserve a safety property in *unbounded time*
- ▶ Use safety proofs to characterize state space:
 - ▶ *Safeable* A SPAS will be available one control action in the future, no matter what action we take at present
 - ▶ *Critical* At least one SPAS may be initiated in the present, but there may not be one available in the future
 - ▶ *Unsafe* No SPAS is available – some sequences may result in safe operation but outcome is uncertain
 - ▶ *Fail* No SPAS available – any action sequence results in guaranteed violation of safety property

STATE SPACE



$S_{\Lambda(\Sigma)}(s_0, \lambda, t)$ is the system's state at time t during its operation.
 $\lambda \in \Lambda(\Sigma)$ is a realization from the probability space of
non-deterministic trajectories available when the system
satisfies the constraints of scenario Σ .

SAFETY PREDICATES

- ▶ A set of safety objectives $\{P_1(s, \Pi_1), P_2(s, \Pi_2), \dots\}$ along with a corresponding set of safety predicates for each, $\{\Psi_1(s_0, \Sigma, \Pi_1), \Psi_2(s_0, \Sigma, \Pi_2), \dots\}$
 - ▶ For objective i , $\Psi_i(s_0, \Sigma, \Pi_i)$ evaluates safety from the present state s_0 into the future in unbounded time
 - ▶ Account for a range of future possibilities associated with a sequence of future actions and uncertain system dynamics
- ▶ Safety predicates Ψ_i are quantifier-free, ensure safety for scenario Σ
- ▶ Scenario Σ represents future uncertainty and action sequences for multiple actors

PROOF REQUIREMENT: GUARANTEEING SAFETY PROPERTY IN UNBOUNDED TIME

Theorem *In-model safety in unbounded time*

$$\begin{aligned} \forall s_0, \Sigma, \Pi_i, \Psi_i(s_0, \Sigma, \Pi_i) \rightarrow \\ \forall t, \lambda \in \Lambda(\Sigma), t > 0 \rightarrow P_i(S_{\Lambda(\Sigma)}(s_0, \lambda, t), \Pi_i) \wedge \\ \Psi_i(S_{\Lambda(\Sigma)}(s_0, \lambda, t), \Sigma \triangleright t, \Pi_i) \quad (1) \end{aligned}$$

parameterized by initial state s_0 , scenario Σ , and safety parameters Π_i .

This says each Ψ_i evaluates safety property for scenario Σ in unbounded time under model assumptions.

HOW TO HANDLE OFF-MODEL DYNAMICS

- ▶ What happens when the system dynamics violate our proof assumptions?
 - ▶ We can call this off-model dynamics
- ▶ Safety properties may not hold
- ▶ Safety controller must not “give up”
- ▶ Still give safest advice it can under the circumstances
- ▶ We need to quantify a spectrum of weaker safety guarantees when the strong property fails

PROOF REQUIREMENT: RANKING ACTION SEQUENCES FOR SAFETY

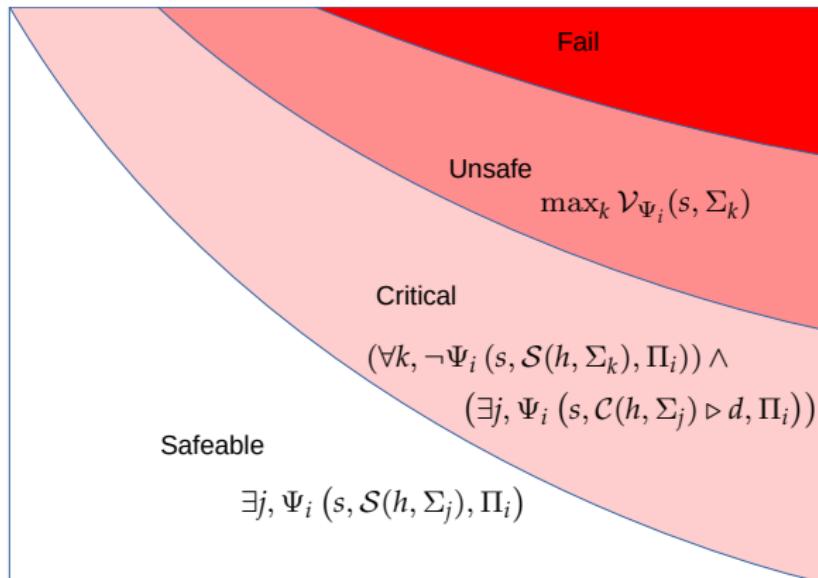
To handle off-model – unplanned, unmodeled – system dynamics, we define the set of safe parameter possibilities $\mathcal{R}_{\Psi_i}(s, \Sigma) = \{\alpha \mid \Psi_i(s, \Sigma, \alpha)\}$, and require a monotonic function in parameter space $\mathcal{V}_{\Psi_i}(s, \Sigma)$.

Theorem *Off-model monotonic safety ranking*

$$\forall s, \Sigma, \Xi, \mathcal{V}_{\Psi_i}(s, \Sigma) < \mathcal{V}_{\Psi_i}(s, \Xi) \rightarrow \mathcal{R}_{\Psi_i}(s, \Sigma) \subset \mathcal{R}_{\Psi_i}(s, \Xi) \quad (2)$$

We can use $\mathcal{V}_{\Psi_i}(s, \Sigma)$ as a metric that helps us quantify an action plan's relative "closeness" to preserving the safety property.

STATE SPACE



GUARANTEES

- ▶ Each control action we check for safety-preserving action sequences available for future control actions
- ▶ When there is no safety-preserving action sequence in the future, one is guaranteed available in the present due to analysis from the previous timestep
- ▶ Safeable states cannot transition directly to unsafe states without first transitioning through a critical state (guaranteed by formal proof)

OUTLINE

- ▶ Introduction
- ▶ Approach
- ▶ **Controller Synthesis**
- ▶ Implementation and Results
- ▶ Conclusions

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
              1:n), or, 1:k) then
    (an, d) = (ap, zeros(size(d)))
    h = (h ⊲ T) ++ (Σ₀ ⊲ T)
  else
    a = map(λi.λe.ζ(Σi, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap, a, v, m)
    end
    d = map(λi.λe.(an==ζ(Σi, e)) ? e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h ⊲ T) ++ (Σn ▷ d[n] ⊲ T)
  end
  return an
end
end
```

- We can synthesize a controller that is generally applicable to systems that have safety predicates as described
- The function body is invoked at each control action, with up-to-date state information and the action proposed by the primary controller at this moment

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψⱼ(s, S(h, Σᵢ), Πⱼ), and,
              1:n), or, 1:k) then
    (an, d) = (ap, zeros(size(d)))
    h = (h▷T) ++ (Σ₀ ⊲ T)
  else
    a = map(λi.λe.ζ(Σᵢ, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψⱼ(s, C(h, Σᵢ ▷ e), Πⱼ), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψⱼ(s, C(h, Σᵢ ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap, a, v, m)
    end
    d = map(λi.λe.(an==ζ(Σᵢ, e))?e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h▷T) ++ (Σₙ▷d[n] ⊲ T)
  end
  return an
end
end
```

} Compute safety predicates to identify
safeable states: is there any safety-
preserving action sequence that we
could implement starting in the future?

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
    if mapreduce(λi.
        mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
        1:n, or, 1:k) then
        (an, d) = (ap, zeros(size(d)))
        h = (h▷T) ++ (Σ₀ ⊲ T)
    else
        a = map(λi.λe.ζ(Σi, e), 1:k, d)
        v = map(λi.λe.
            map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
            1:k, d)
        m = map(λi.λe.
            map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
            1:k, d)
        if mapreduce(λj.reduce(and, v[j]), or,
                    findindices(λq.q==ap, a))
            an = ap
        else
            an = U(ap, a, v, m)
        end
        d = map(λi.λe.(an==ζ(Σi, e))?e+T:0,
                1:k, d)
        n = findindices(λq.q==an, a)[1]
        h = (h▷T) ++ (Σn ▷ d[n] ⊲ T)
    end
    return an
end
end
```

} If so, allow primary controller to act
without interference.

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
              1:n), or, 1:k) then
    (an, d) = (ap, zeros(size(d)))
    h = (h▷T) ++ (Σ₀ ⊲ T)
  else
    a = map(λi.λe.ζ(Σi, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap, a, v, m)
    end
    d = map(λi.λe.(an==ζ(Σi, e))?e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h▷T) ++ (Σn ▷ d[n] ⊲ T)
  end
  return an
end
end
```

} Compute next action in sequence under
a range of future decisions.

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψⱼ(s, S(h, Σᵢ), Πⱼ), and,
              1:n), or, 1:k) then
    (an,d) = (ap,zeros(size(d)))
    h = (h▷T)++(Σ₀⊲T)
  else
    a = map(λi.λe.ζ(Σᵢ, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψⱼ(s, C(h, Σᵢ ▷ e), Πⱼ), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψⱼ(s, C(h, Σᵢ ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap, a, v, m)
    end
    d = map(λi.λe.(an==ζ(Σᵢ, e))?e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h▷T)++(Σₙ▷d[n]⊲T)
  end
  return an
end
end
```

Compute safety predicates to identify critical states: which safety-preserving action sequence can we implement starting now? which is underway that we can continue?

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
    if mapreduce(λi.
        mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
        1:n), or, 1:k) then
        (an,d) = (ap,zeros(size(d)))
        h = (h▷T)++(Σ₀⊲T)
    else
        a = map(λi.λe.ζ(Σi, e), 1:k, d)
        v = map(λi.λe.
            map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
            1:k, d)
        m = map(λi.λe.
            map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
            1:k, d)
        if mapreduce(λj.reduce(and, v[j]), or,
                    findindices(λq.q==ap, a))
            an = ap
        else
            an = U(ap, a, v, m)
        end
        d = map(λi.λe.(an==ζ(Σi, e))?e+T:0,
                1:k, d)
        n = findindices(λq.q==an, a)[1]
        h = (h▷T)++(Σn▷d[n]⊲T)
    end
    return an
end
end
```

} Compute ranking to evaluate safety
characteristics of action sequences.

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
              1:n), or, 1:k) then
    (an,d) = (ap,zeros(size(d)))
    h = (h▷T)++(Σ₀⊲T)
  else
    a = map(λi.λe.ζ(Σi, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap,a,v,m)
    end
    d = map(λi.λe.(an==ζ(Σi, e))?e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h▷T)++(Σn▷d[n]⊲T)
  end
  return an
end
end
```

} Invoke action that implements or continues a safety-preserving action sequence.

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
              1:n), or, 1:k) then
    (an,d) = (ap,zeros(size(d)))
    h = (h▷T)++(Σ₀⊲T)
  else
    a = map(λi.λe.ζ(Σi, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap, a, v, m)
    end
    d = map(λi.λe.(an==ζ(Σi, e))?e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h▷T)++(Σn▷d[n]⊲T)
  end
  return an
end
end
```

If there is no available safety-preserving action sequence, which is } most safe or maintains the most desirable combination safety properties (as determined by function U)?

CONTROLLER SYNTHESIS

```
let d = zeros(k), h = Σ₀ ⊲ g
function control_action(s, ap)
  if mapreduce(λi.
    mapreduce(λj.Ψj(s, S(h, Σi), Πj), and,
              1:n), or, 1:k) then
    (an,d) = (ap,zeros(size(d)))
    h = (h▷T)++(Σ₀▷T)
  else
    a = map(λi.λe.ζ(Σi, e), 1:k, d)
    v = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e), Πj), 1:n),
      1:k, d)
    m = map(λi.λe.
      map(λj.Ψj(s, C(h, Σi ▷ e)), 1:n),
      1:k, d)
    if mapreduce(λj.reduce(and, v[j]), or,
                findindices(λq.q==ap, a))
      an = ap
    else
      an = U(ap,a,v,m)
    end
    d = map(λi.λe.(an==ζ(Σi, e))?e+T:0,
            1:k, d)
    n = findindices(λq.q==an, a)[1]
    h = (h▷T)++(Σn▷d[n]▷T)
  end
  return an
end
end
```

} History tracks system action sequence
that has been invoked but due to sys-
tem lag not yet affected system dynam-
ics (but will in the future)

OUTLINE

- ▶ Introduction
- ▶ Approach
- ▶ Controller Synthesis
- ▶ **Implementation and Results**
- ▶ Conclusions

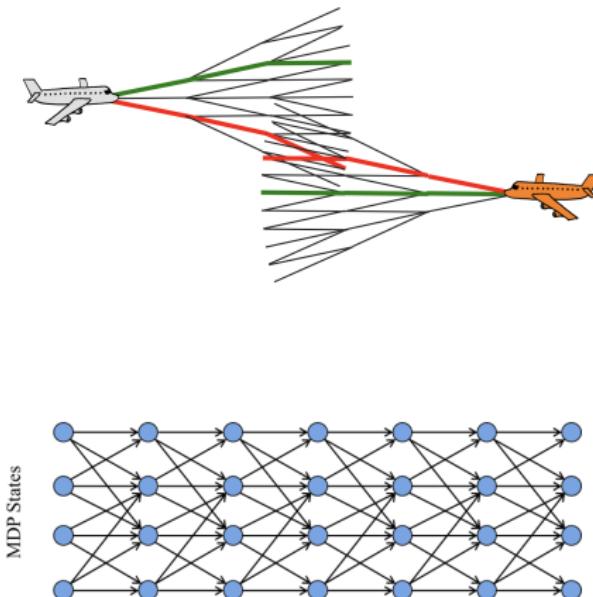
AIRBORNE COLLISION AVOIDANCE

- ▶ Industrial systems install on all airliners
 - ▶ Detects potential mid-air collisions
 - ▶ Since 1980s: TCAS
 - ▶ New system: ACAS-X
- ▶ Give vertical advisories, e.g. *Climb, Do not descend*
- ▶ System of last resort



ACAS-X DESIGN

- ▶ Core of the system is a table driving behavior
 - ▶ Optimal policy for a Markov Decision Process (discrete dynamics)
 - ▶ Cost for each action at discrete cutpoints in the state space
 - ▶ Large size (36 MB compressed, 2GB uncompressed)
- ▶ Additional online costs computed to modify costs according to ad-hoc rules



SAFETY ANALYSIS

We developed predicates for safety analysis of vertical collision avoidance maneuvers in ACAS X:

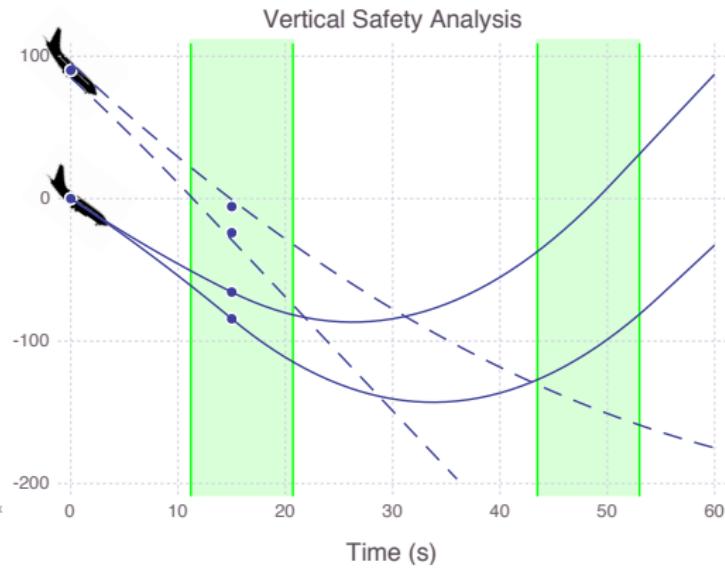
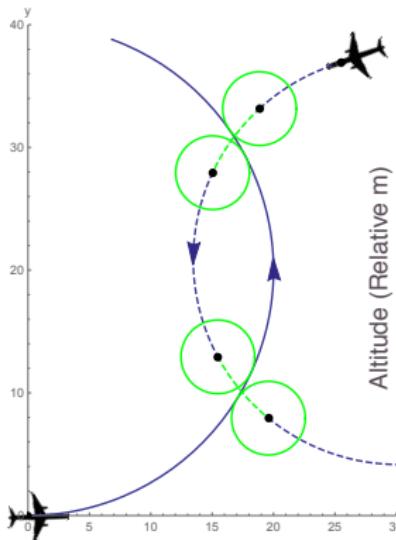
Kouskoulas, et al., *Formally Verified Safe Vertical Maneuvers for Non-Deterministic, Accelerating Aircraft Dynamics*, In Interactive Theorem Proving (ITP) , 2017.

- ▶ Predicates used to analyze an existing system's safety properties, not synthesize a controller

SAFETY ANALYSIS

- ▶ Vertical predicates used for safety analysis of ACAS X
 - ▶ Use an independent model for analyzing safety in an existing system
 - ▶ Developed using formal methods (FM), rigorous mathematical approaches to modeling and proving properties about a software system
 - ▶ Use hybrid system model combining discrete, pilot decisions and sequences of discrete actions with continuous dynamics
 - ▶ Use higher-order logic (HOL) and calculus of inductive constructions (CIC) to state safety conjectures in continuous time (Coq theorem prover)
 - ▶ With minor changes, predicates satisfy requirements for provably correct controller synthesis described in our approach.

EXAMPLE VERTICAL SAFETY ANALYSIS



VERTICAL SAFETY PREDICATE

$$\begin{aligned}
\Gamma((A, B, C), t_b, t_e) \equiv t_b \leq t_e \rightarrow \\
(A > 0 \wedge ((0 \leq D \wedge (R_1 > t_e \vee R_2 < t_b)) \vee D < 0) \vee \\
A < 0 \wedge (0 < D \wedge R_2 < t_b \wedge R_1 > t_e) \vee \\
A = 0 \wedge (B > 0 \wedge -C/B < t_b \vee B < 0 \wedge -C/B > t_e \vee B = 0 \wedge C > 0)) \quad (1)
\end{aligned}$$

$$\begin{aligned}
\Phi(Q_1, L_1, t_{l1}, Q_2, L_2, t_{l2}, t_b, t_e) \equiv \\
\Gamma(Q_1 - Q_2 - P, \max(t_b, 0), \min(t_e, t_{l1}, t_{l2})) \wedge \Gamma(L_1 - L_2 - P, \max(t_b, t_{l1}, t_{l2}), t_e) \wedge \\
(t_{l1} > t_{l2} \rightarrow \Gamma(Q_1 - L_2 - P, \max(t_b, \min(t_{l1}, t_{l2})), \min(t_e, \max(t_{l1}, t_{l2})))) \wedge \\
(t_{l1} < t_{l2} \rightarrow \Gamma(L_1 - Q_2 - P, \max(t_b, \min(t_{l1}, t_{l2})), \min(t_e, \max(t_{l1}, t_{l2})))) \quad (2)
\end{aligned}$$

$$\Psi = \bigwedge_{j \in \{1, \dots, n\}} \left(\left(\bigwedge_{i \in \{1, \dots, m\}} \Phi(Q_i^{\text{Own}\downarrow}, L_i^{\text{Own}\downarrow}, t_{ri}^{\text{Own}\downarrow}, Q_i^{\text{Int}\uparrow}, L_i^{\text{Int}\uparrow}, t_{ri}^{\text{Int}\uparrow}, \tau_{ij}, v_{ij}) \right) \vee \right. \\
\left. \left(\bigwedge_{i \in \{1, \dots, m\}} \Phi(Q_i^{\text{Int}\downarrow}, L_i^{\text{Int}\downarrow}, t_{ri}^{\text{Int}\downarrow}, Q_i^{\text{Own}\uparrow}, L_i^{\text{Own}\uparrow}, t_{ri}^{\text{Own}\uparrow}, \tau_{ij}, v_{ij}) \right) \right) \quad (3)$$

RESULTS

- ▶ Comprehensive safety proof is not possible for ACAS X under our model assumptions
- ▶ Disproved global safety by generating counterexamples
 - ▶ A counterexample is a point in the state space where guaranteed safe advice is possible, but the system does not issue it
- ▶ ACAS X trades off safety under our acceleration assumptions for particular operational behavior
 - ▶ Excessive alerting may lead to unresponsive pilots

SAFETY CONTROLLER SYNTHESIS CASE STUDY

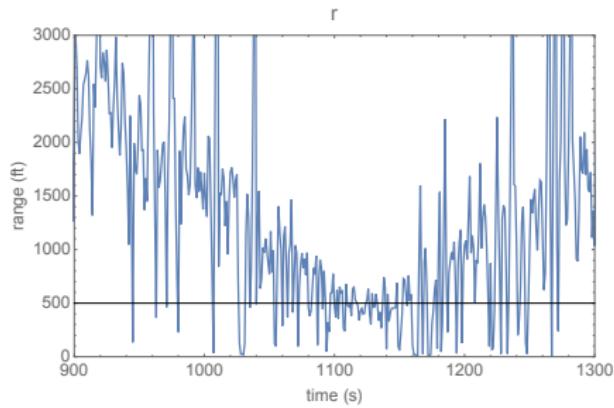
- ▶ It may be that dangerous collisions that need correction only exist in a fraction of a percent of the state space – critical area is important to handle correctly
- ▶ Evaluate ACAS X system performance improvement with additional safety controller
 - ▶ Create pathologically dangerous encounters, chosen specifically in critical parts of the state space where ACAS X gives advice that allows a collision under model assumptions
 - ▶ ACAS X could issue advice that definitively avoid collisions under the model assumptions, but chooses alternative advice
 - ▶ Modified proofs and predicates used for safety analysis to fit the framework described in this paper
 - ▶ Formal, mathematical guarantees of correctness
 - ▶ Synthesized controlled using safety proofs

IN-MODEL PERFORMANCE

- In simulation, the safety controller intervened and eliminated all mid-air collisions resulting in fully safe operation under our dynamics assumptions

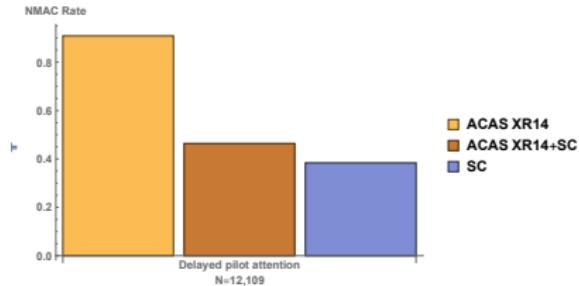
OFF-MODEL PERFORMANCE

- ▶ What happens to controller performance when the dynamics model is violated?
 - ▶ Mathematical analysis for a CPS relies upon model
- ▶ Introduce uncharacterized, off-model noise

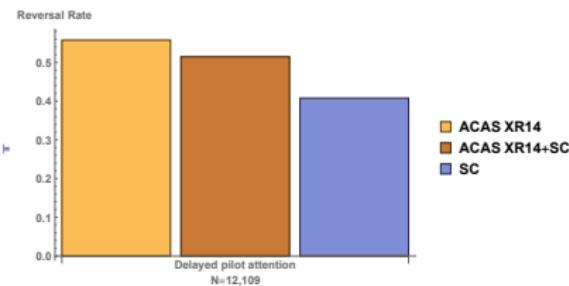


Horizontal range during a two-aircraft encounter output by the ACAS X Run14 surveillance and tracking module over time.
True range would be a smooth curve.

OFF-MODEL PERFORMANCE



Safety performance on a set of stressing encounters passing through from critical states comparing ACAS X to various configurations of the safety controller (SC).



Advisory reversals on a set of stressing encounters passing through critical states comparing ACAS X to various configurations of the safety controller (SC).

OUTLINE

- ▶ Introduction
- ▶ Approach
- ▶ Controller Synthesis
- ▶ Implementation and Results
- ▶ **Conclusions**

CONCLUSIONS FOR ACAS X CASE STUDY

- ▶ Approach produces a practical safety controller for ACAS X that shows good performance when handling uncharacterized noise
- ▶ It does not adversely affect ACAS X in other parts of the state space
- ▶ Encounters sampled by regular testing of ACAS X do not contain a statistically significant representation of state-space critical conditions

GENERAL CONCLUSIONS

- ▶ Promising approach to producing safety controller for complex behavior in CPS systems
 - ▶ Works with a finite dictionary of safety-preserving action sequences
 - ▶ Allow primary controller freedom to pursue objectives until safety cannot be guaranteed
 - ▶ Safety-preserving action sequence always available under model assumptions
 - ▶ Control actions during off-model system evolution chosen to maximize safety insofar as possible, given future uncertainty
 - ▶ Independent of primary controller algorithm, so safety controller doesn't depend on its behavior